

**CSCI235/CSCI835 Database Systems**  
**Assignment 1**  
15 March 2020

---

**Scope**

This assignment includes the tasks related to indexing of relational tables, implementation of data retrieval in PL/SQL and implementation of stored procedures and functions in PL/SQL.

The outcomes of the laboratory work are due by **Saturday 18 April, 2020, 7.00 pm (sharp)**.

**Please read very carefully information listed below.**

This laboratory contributes to 10% of the total evaluation in a subject CSCI235 and it contributes to 6% of the total evaluation in a subject CSCI835.

A submission procedure is explained at the end of specification.

This assignment consists of 4 tasks and specification of each task starts from a new page.

It is recommended to solve the problems before attending the laboratory classes in order to efficiently use supervised laboratory time.

A submission marked by Moodle as "late" is treated as a late submission no matter how many seconds it is late.

A policy regarding late submissions is included in the subject outline.

A submission of compressed files (zipped, gzipped, rared, tared, 7-zipped, lhzed, ... etc) is not allowed. The compressed files will not be evaluated.

All files left on Moodle in a state "Draft (not submitted) " will not be evaluated.

An implementation that does not compile due to one or more syntactical errors scores no marks.

It is expected that all tasks included within **Assignment 1** will be solved **individually without any cooperation** with the other students. If you have any doubts, questions, etc. please consult your lecturer or tutor during lab classes or office hours. Plagiarism will result in a **FAIL** grade being recorded for the assessment task.

---

## Task 1 (2 marks)

### Prologue

Download the files `dbcreate.sql` and `dbdrop.sql` included in a section **SAMPLE DATABASE**. To drop a sample database, process a script `dbdrop.sql`. To create a sample database, process as script `dbcreate.sql`. It is strongly recommended to drop a sample database and to re-create it before implementation of each task.

Connect to Oracle database server and process the following SQL statement that saves a query processing plan for a given `SELECT` statement in `PLAN_TABLE`.

```
EXPLAIN PLAN FOR SELECT ORDER_ID, ORDER_DATE FROM ORDERS;
```

Next, process the following `SELECT` statement to display a query processing plan stored in `PLAN_TABLE`.

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

Among the others, you should get the following results.

PLAN\_TABLE\_OUTPUT

-----  
Plan hash value: 1275100350

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		131	2882	3 (0)	00:00:01
1	TABLE ACCESS FULL	ORDERS	131	2882	3 (0)	00:00:01

A line **TABLE ACCESS FULL| ORDERS** in a plan given above indicates that a database system plans to access a table `ORDERS` to compute the query.

Next, create an index on the columns `ORDER_ID` and `ORDER_DATE` in a relational table `ORDERS`.

```
CREATE INDEX ORDERS_IDX ON ORDERS(ORDER_ID, ORDER_DATE);
```

Again, process the following SQL statements that save a query processing plan for the same `SELECT` statement as before in `PLAN_TABLE` and display a query processing plan stored in `PLAN_TABLE`.

```
EXPLAIN PLAN FOR SELECT ORDER_ID, ORDER_DATE FROM ORDERS;
```

Among the others, you should get the following results.

PLAN\_TABLE\_OUTPUT

-----  
Plan hash value: 2467194144

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		131	2882	1 (0)	00:00:01
1	INDEX FULL SCAN	ORDERS_IDX	131	2882	1 (0)	00:00:01

This time a database system plans to use an index `ORDERS_IDX` created a moment ago to process the same query. Note, a line `INDEX FULL SCAN | ORDERS_IDX` in a plan given above means that a database system plans to horizontally traverse leaf level an index `ORDERS_IDX` to find the values in the columns `ORDER_ID` and `ORDER_DATE`.

### Conclusions

`EXPLAIN PLAN` statement of SQL can be used to get information about a processing plan created by a query processor for a given `SELECT` statement. A query processing plan provides information on whether and index created earlier will be used for processing of SQL statement. We shall use `EXPLAIN PLAN` statement to check whether an index created to speed up `SELECT` statement will be used for processing of the statement.

To drop an index, process a statement

```
DROP INDEX ORDERS_IDX;
```

No report is expected from processing of SQL statements given above.

### Problem

Your task is to find what indexes should be created to speed up processing of `SELECT` statements listed below. You are expected to create one index for one `SELECT` statement. To simplify the problem, assume that any index which is later on used by a query processor to speed up processing of `SELECT` statement will do.

- (i) 

```
SELECT *  
FROM ORDER_DETAIL  
WHERE PRODUCT_NAME = 'BOLT' AND  
       QUANTITY > 100;
```
- (ii) 

```
SELECT DISTINCT CATEGORY_NAME  
FROM PRODUCT;
```

- (iii) 

```
SELECT UNIT_PRICE
FROM ORDER_DETAIL
WHERE QUANTITY IN (100, 200, 300) OR
      DISCOUNT = 0.01;
```
- (iv) 

```
SELECT CATEGORY_NAME, SUPPLIER_NAME, COUNT(*)
FROM PRODUCT
GROUP BY CATEGORY_NAME, SUPPLIER_NAME;
```
- (v) 

```
SELECT SUPPLIER_NAME, UNIT_PRICE
FROM PRODUCT
ORDER BY UNIT_PRICE, QUANTITY_PER_UNIT;
```

Implement SQL script `solution1.sql` such that for each one of SELECT statements given above the script performs the following actions.

- (i) Find and list a query processing plan for SELECT statement without an index.
- (ii) Create an index.
- (iii) Find and list a query processing plan for SELECT statement with an index.
- (iv) Drop an index.

When ready process SQL script file `solution1.sql` and save a report from processing in a file `solution1.lst`.

Your report must include a listing of all PL/SQL statements processed. To achieve that put the following SQLcl commands:

```
SPOOL solution1
SET ECHO ON
SET FEEDBACK ON
SET LINESIZE 300
SET PAGESIZE 200
```

at the beginning of SQL script and

```
SPOOL OFF
```

at the end of SQL script.

### **Deliverables**

A file `solution1.lst` with a report from processing of a script file `solution1.sql` that lists query processing plans before and after indexing. A report must have no errors and it must list all SQL statements processed.

---

**Task 2 (3 marks)**

Implement an anonymous PL/SQL block that lists order id (attribute `ORDER_ID`), company name that submitted order (attribute `COMPANY_NAME`), and order date (attribute `ORDER_DATE`) from the five most recently submitted orders.

To list information retrieved from a sample database use PL/SQL package `DBMS_OUTPUT`. It is explained in the Cookbook, Recipe 7.1 How to start programming in PL/SQL how to use `DBMS_OUTPUT` package. Remember about `SET SERVEROUTPUT ON` at the beginning of a script file that contains your anonymous PL/SQL block.

Information retrieved must be listed in the following format.

```
Order id: 777
Order date: 26-APR-98
Company name: Golden Bolts Pty Ltd
=====
Order id: 666
Order date: 11-MAR-96
Company name: Lazy Lobster Seafood Corp
=====
...
and so on.
```

Your implementation must use at least one cursor and at least one exception handler. In fact, such constraints make your implementation easier.

To test your solution put an implemented anonymous PL/SQL block into SQL script file `solution2.sql` and process the script.

Your report must include a listing of all PL/SQL statements processed. To achieve that put the following SQLcl commands:

```
SPOOL solution2
SET SERVEROUTPUT ON
SET ECHO ON
SET FEEDBACK ON
SET LINESIZE 100
SET PAGESIZE 200
SET SERVEROUTPUT ON
```

at the beginning of SQL script and

```
SPOOL OFF
```

at the end of SQL script.

**Deliverables**

A file `solution2.lst` with a report from testing of an anonymous PL/SQL block implemented in this task. A report must have no errors and it must list all PL/SQL and SQL statements processed.

---

**Task 3 (3 marks)**

Implement a stored PL/SQL procedure

```
INSERT_ORDER_DETAIL (order_id, product_name, unit_price, quantity, discount)
```

that inserts a row into a relational table ORDER\_DETAIL and enforces the following consistency constraint on data entry into a relational table ORDER\_DETAIL.

*A product can be ordered only if it is not discontinued.*

If the consistency constraint is satisfied insert and commit a row in ORDER\_DETAIL table. Otherwise, use DBMS\_OUTPUT PL/SQL package to display an error message when the consistency constraint is violated and do not insert a row.

When INSERT\_ORDER\_DETAIL procedure is ready create SQL script solution3.sql that stores the procedure in a data dictionary and tests the procedure with two EXECUTE statements. First, test the procedure for a product that is not discontinued and then test it again for a product that is discontinued. Any discontinued and not discontinued products used for testing will do.

Process SQL script solution3.sql and save a report from processing in a file solution3.lst.

Your report must include a listing of all PL/SQL statements processed. To achieve that put the following SQLcl commands:

```
SPOOL solution3
SET SERVEROUTPUT ON
SET ECHO ON
SET FEEDBACK ON
SET LINESIZE 100
SET PAGESIZE 200
SET SERVEROUTPUT ON
```

at the beginning of SQL script and

```
SPOOL OFF
```

at the end of SQL script.

**Deliverables**

A file solution3.lst with a report from processing of SQL script solution3.sql. A report must have no errors and it must list all PL/SQL and SQL statements processed.

---

**Task 4 (2 marks)**

Implement a stored PL/SQL function

```
TOTAL_ORDERS (company_name)
```

that returns the total number of orders submitted by a given customer. Make a company name a parameter of a stored function. Assume that company name uniquely identifies each customer.

When ready, implement a script `solution4.sql` that stores the function in a data dictionary and tests a function. To test a function implement `SELECT` statements that lists company name (attribute `COMPANY_NAME`), address (attribute `ADDRESS`), and the total number of submitted orders for all companies that submitted more than 2 and less than 6 orders.

Process SQL script `solution4.sql` and save a report from processing in a file `solution4.lst`.

Your report must include a listing of all PL/SQL statements processed. To achieve that put the following SQLcl commands:

```
SPOOL solution4
SET SERVEROUTPUT ON
SET ECHO ON
SET FEEDBACK ON
SET LINESIZE 100
SET PAGESIZE 200
SET SERVEROUTPUT ON
```

at the beginning of SQL script and

```
SPOOL OFF
```

at the end of SQL script.

**Deliverables**

A file `solution4.lst` with a report from processing of SQL script `solution4.sql`. A report must have no errors and it must list all PL/SQL and SQL statements processed.

---



### **Submission**

Submit the files **solution1.lst**, **solution2.lst**, **solution3.lst**, and **solution4.lst** through Moodle in the following way:

- (1) Access Moodle at **<http://moodle.uowplatform.edu.au/>**
- (2) To login use a **Login** link located in the right upper corner the Web page or in the middle of the bottom of the Web page
- (3) When logged select a site **CSCI835/CSCI235 (S120) Database Systems**
- (4) Scroll down to a section **SUBMISSIONS**
- (5) Click at a link **In this place you can submit the outcomes of Assignment 1**
- (6) Click at a button **Add Submission**
- (7) Move a file **solution1.lst** into an area **You can drag and drop files here to add them**. You can also use a link **Add...**
- (8) Repeat a step (7) for the files **solution2.lst**, **solution3.lst**, and **solution4.lst**.
- (9) Click at a button **Save changes**
- (10) Click at a button **Submit assignment**
- (11) Click at the checkbox with a text attached: **By checking this box, I confirm that this submission is my own work, ...** in order to confirm the authorship of your submission.
- (12) Click at a button **Continue**

---

*End of specification*