

Table of Contents

Original cost of query:.....	2
Optimization #1 – Transformation of Select Statement.....	3
Optimization #2 – Creation of Bitmap Index	6
Optimization #3 – Creation of B*Tree Index	9
Optimization #4 – Allocation of Additional db_cache_size	12
utlbstat/utlestat (Screenshots used because of spacing issues)	14
Total Costs.....	16

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

Original cost of query:

	Id		Operation		Name		Rows		Bytes		TempSpc		Cost	(%CPU)		Time	
	0		SELECT STATEMENT				877K		98M				12273	(1)		00:00:01	
	1		MERGE JOIN				877K		98M				12273	(1)		00:00:01	
	2		SORT JOIN				876K		32M				12179	(1)		00:00:01	
*	3		VIEW		REVENUE		876K		32M				12179	(1)		00:00:01	
	4		WINDOW BUFFER				876K		40M				12179	(1)		00:00:01	
	5		HASH GROUP BY				876K		40M				12179	(1)		00:00:01	

PLAN_TABLE_OUTPUT

*	6		TABLE ACCESS FULL		LINEITEM		876K		40M				12157		(1)		00:00:01	
*	7		SORT JOIN				3046		234K		568K		94		(2)		00:00:01	
	8		TABLE ACCESS FULL		SUPPLIER		3046		234K				34		(0)		00:00:01	

Optimization #1 – Transformation of Select Statement

Note that I decided against using Materialized Views as we're unsure of whether or not the client's DB has designated downtime to mitigate the more expensive insertion, and if the client doesn't have a designated downtime for the DB to perform inserts and commits to update the materialized views and instead wants the views to remain up to date and therefore, updated on insert, if the table LINEITEM is frequently inserted to which is what I assume since it's the largest table, the performance of the entire DB will greatly suffer which is why I've decided not to use Materialized Views.

(1) Description of Improvement:

As seen in the origin plan for the query in row id 3, the view operation is performed indicating that the view was not successfully merged.

I performed a query transformation so that the view will successfully merge.

(2) Benefits of Improvement:

The view now successfully merges with the select query and slightly reduces the cost of the processing plan which can be seen below

```
SQL> SET ECHO ON

SQL> SET FEEDBACK ON

SQL> SET LINESIZE 300

SQL> SET PAGESIZE 300

SQL>

SQL> CREATE VIEW REVENUE( S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, TOTAL_REVENUE ) AS

  2      SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, L_EXTENDEDPRICE * (1 - L_DISCOUNT)
  3      FROM LINEITEM, SUPPLIER
  4      WHERE L_SUPPKEY = S_SUPPKEY
  5      AND LINEITEM.L_SHIPDATE >= '08-NOV-1995'
  6      AND LINEITEM.L_SHIPDATE <= '26-MAY-2002';
```

View REVENUE created.

```
SQL>

SQL> EXPLAIN PLAN FOR

  2  SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
  3  FROM REVENUE
  4  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
  5  HAVING SUM(TOTAL_REVENUE) =
  6      (SELECT MAX(SUM(R.TOTAL_REVENUE))
  7      FROM REVENUE R
  8      GROUP BY R.S_SUPPKEY)
  9  ORDER BY S_SUPPKEY;
```

Explained.

```
SQL>

SQL> @showplan

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 3144199225

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		859K	104M	12215 (1)	00:00:01
* 1	FILTER					
2	SORT GROUP BY		859K	104M	12215 (1)	00:00:01
* 3	HASH JOIN		859K	104M	12194 (1)	00:00:01
4	TABLE ACCESS FULL	SUPPLIER	3109	239K	34 (0)	00:00:01
* 5	TABLE ACCESS FULL	LINEITEM	859K	39M	12157 (1)	00:00:01
6	SORT AGGREGATE		1	48	12179 (1)	00:00:01
7	SORT GROUP BY		1	48	12179 (1)	00:00:01
* 8	TABLE ACCESS FULL	LINEITEM	859K	39M	12157 (1)	00:00:01

Predicate Information (identified by operation id):

```
1 - filter(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))= (SELECT
      MAX(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))) FROM "TPCHR"."LINEITEM"
      "LINEITEM" WHERE "LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08
      00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
      "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss') GROUP BY "L_SUPPKEY"))
3 - access("L_SUPPKEY"="S_SUPPKEY")
5 - filter("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00',
      'syyyy-mm-dd hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE('
      2002-05-26 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND "L_SUPPKEY">=0)
8 - filter("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00',
      'syyyy-mm-dd hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE('
      2002-05-26 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

(3) Costs of Improvement:

There is no cost associated with this improvement.

(4) Report from improvement:

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> CREATE VIEW REVENUE( S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, TOTAL_REVENUE ) AS
2      SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, L_EXTENDEDPRICE * (1 - L_DISCOUNT)
3      FROM LINEITEM, SUPPLIER
4      WHERE L_SUPPKEY = S_SUPPKEY
5      AND LINEITEM.L_SHIPDATE >= '08-NOV-1995'
6      AND LINEITEM.L_SHIPDATE <= '26-MAY-2002';
```

View REVENUE created.

```
SQL>
SQL> SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
2  FROM REVENUE
3  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
4  HAVING SUM(TOTAL_REVENUE) =
5      (SELECT MAX(SUM(R.TOTAL_REVENUE))
6      FROM REVENUE R
7      GROUP BY R.S_SUPPKEY)
8  ORDER BY S_SUPPKEY;
```

S_SUPPKEY	S_NAME	S_ADDRESS	S_PHONE	TOTAL_REVENUE
2993	Supplier#000002993	CDRN7azuEWTawl9G	30-541-514-5637	12267625.5

1 row selected.

Optimization #2 – Creation of Bitmap Index

(1) Description of Improvement:

I created a bitmap index over the L_SHIPDATE column;
the total number of rows in a table LINEITEM is 1,800,093
the total number of distinct rows of a column L_SHIPDATE is 2,526
meaning that the cardinality of L_SHIPDATE = $(2,526/1,800,093) * 100 = 0.14\%$ making the L_SHIPDATE a good column for a bitmap to be created.

```
CREATE BITMAP INDEX TASK1IDX1 ON LINEITEM(L_SHIPDATE);
```

(2) Benefits of Improvement:

The bitmap index significantly improves the performance of the query since the LINEITEM table is so large.
Rather than having to traverse the entire LINEITEM table, the query optimizer can instead traverse the bitmap index.

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> CREATE BITMAP INDEX TASK1IDX1 ON LINEITEM(L_SHIPDATE);

INDEX TASK1IDX1 created.
SQL>
SQL> EXPLAIN PLAN FOR
  2  SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
  3  FROM REVENUE
  4  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
  5  HAVING SUM(TOTAL_REVENUE) =
  6      (SELECT MAX(SUM(R.TOTAL_REVENUE))
  7      FROM REVENUE R
  8      GROUP BY R.S_SUPPKEY)
  9  ORDER BY S_SUPPKEY;
```

Explained.

```
SQL>
SQL> @showplan
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 3552806918

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		859K	104M	1406 (2)	00:00:01
* 1	FILTER					
2	SORT GROUP BY		859K	104M	1406 (2)	00:00:01
* 3	HASH JOIN		859K	104M	1385 (1)	00:00:01
4	TABLE ACCESS FULL	SUPPLIER	3109	239K	34 (0)	00:00:01

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

*	5		TABLE ACCESS BY INDEX ROWID BATCHED		LINEITEM		859K		39M		1348	(0)		00:00:01	
	6		BITMAP CONVERSION TO ROWIDS												
*	7		BITMAP INDEX RANGE SCAN		TASK1IDX1										
	8		SORT AGGREGATE				1		48		1370	(2)		00:00:01	
	9		SORT GROUP BY				1		48		1370	(2)		00:00:01	
	10		TABLE ACCESS BY INDEX ROWID BATCHED		LINEITEM		859K		39M		1348	(0)		00:00:01	
	11		BITMAP CONVERSION TO ROWIDS												
*	12		BITMAP INDEX RANGE SCAN		TASK1IDX1										

Predicate Information (identified by operation id):

```
1 - filter(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))= (SELECT
      MAX(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))) FROM "TPCHR"."LINEITEM" "LINEITEM" WHERE
      "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
      "LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd hh24:mi:ss')
GROUP BY
      "L_SUPPKEY"))
3 - access("L_SUPPKEY"="S_SUPPKEY")
5 - filter("L_SUPPKEY">=0)
7 - access("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss'))
12 - access("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss'))
```

(3) Costs of Improvement:

The cost of creating the index is 6.5MB in persistent storage.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='&INDEX_NAME';

old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='&INDEX_NAME'

new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='TASK1IDX1'
```

Index Size (MB)

6.5

(4) Report from improvement:

```
SQL> SET ECHO ON

SQL> SET FEEDBACK ON

SQL> SET LINESIZE 300

SQL> SET PAGESIZE 300

SQL>

SQL> CREATE BITMAP INDEX TASK1IDX1 ON LINEITEM(L_SHIPDATE);

INDEX TASK1IDX1 created.
```

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

```
SQL>
SQL> SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
2  FROM REVENUE
3  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
4  HAVING SUM(TOTAL_REVENUE) =
5      (SELECT MAX(SUM(R.TOTAL_REVENUE))
6      FROM REVENUE R
7      GROUP BY R.S_SUPPKEY)
8  ORDER BY S_SUPPKEY;
```

S_SUPPKEY	S_NAME	S_ADDRESS	S_PHONE	TOTAL_REVENUE
2993	Supplier#000002993	CDRN7azuEWTaw19G	30-541-514-5637	12267625.5

1 row selected.

Optimization #3 – Creation of B*Tree Index

(1) Description of Improvement:

I created a B*Tree index over the S_SUPPKEY, S_NAME, S_ADDRESS and S_PHONE columns.

```
CREATE INDEX TASK1IDX2 ON SUPPLIER(S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE);
```

(2) Benefits of Improvement:

Rather than accessing the entire SUPPLIER table, the query optimizer can traverse the index vertically.

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> CREATE INDEX TASK1IDX2 ON SUPPLIER(S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE);

Index TASK1IDX2 created.
SQL>
SQL> EXPLAIN PLAN FOR
  2  SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
  3  FROM REVENUE
  4  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
  5  HAVING SUM(TOTAL_REVENUE) =
  6      (SELECT MAX(SUM(R.TOTAL_REVENUE))
  7      FROM REVENUE R
  8      GROUP BY R.S_SUPPKEY)
  9  ORDER BY S_SUPPKEY;
```

Explained.

```
SQL>
SQL> @showplan
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 853867656

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		859K	104M	1390 (2)	00:00:01
* 1	FILTER					
2	SORT GROUP BY		859K	104M	1390 (2)	00:00:01
* 3	HASH JOIN		859K	104M	1369 (1)	00:00:01
4	INDEX FAST FULL SCAN	TASK1IDX2	3109	239K	18 (0)	00:00:01
* 5	TABLE ACCESS BY INDEX ROWID BATCHED	LINEITEM	859K	39M	1348 (0)	00:00:01
6	BITMAP CONVERSION TO ROWIDS					
* 7	BITMAP INDEX RANGE SCAN	TASK1IDX1				

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

	8		SORT AGGREGATE				1		48		1370	(2)		00:00:01	
	9		SORT GROUP BY				1		48		1370	(2)		00:00:01	
	10		TABLE ACCESS BY INDEX ROWID BATCHED		LINEITEM		859K		39M		1348	(0)		00:00:01	
	11		BITMAP CONVERSION TO ROWIDS												
	* 12		BITMAP INDEX RANGE SCAN		TASK1IDX1										

Predicate Information (identified by operation id):

```
1 - filter(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))= (SELECT
      MAX(SUM("L_EXTENDEDPRICE"*(1-"L_DISCOUNT"))) FROM "TPCHR"."LINEITEM" "LINEITEM" WHERE
      "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd hh24:mi:ss') AND
      "LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd hh24:mi:ss')
GROUP BY
      "L_SUPPKEY"))
3 - access("L_SUPPKEY"="S_SUPPKEY")
5 - filter("L_SUPPKEY">=0)
7 - access("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss'))
12 - access("LINEITEM"."L_SHIPDATE">=TO_DATE(' 1995-11-08 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss') AND "LINEITEM"."L_SHIPDATE"<=TO_DATE(' 2002-05-26 00:00:00', 'syyyy-mm-dd
      hh24:mi:ss'))
```

(3) Costs of Improvement:

The cost of creating the index is 0.5MB in persistent storage.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='&INDEX_NAME';

old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='&INDEX_NAME'

new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where
segment_name='TASK1IDX2'
```

Index Size (MB)

 .5

(4) Report from improvement:

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> CREATE INDEX TASK1IDX2 ON SUPPLIER(S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE);
```

Index TASK1IDX2 created.

```
SQL>
SQL> SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
2  FROM REVENUE
3  GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
4  HAVING SUM(TOTAL_REVENUE) =
5      (SELECT MAX(SUM(R.TOTAL_REVENUE))
6      FROM REVENUE R
7      GROUP BY R.S_SUPPKEY)
8  ORDER BY S_SUPPKEY;
```

S_SUPPKEY	S_NAME	S_ADDRESS	S_PHONE	TOTAL_REVENUE
2993	Supplier#000002993	CDRN7azuEWTawl9G	30-541-514-5637	12267625.5

1 row selected.

Optimization #4 – Allocation of Additional db_cache_size

Note that this optimization was completed after I had completed my optimizations for all 5 tasks. The other tasks were tested for read operation and query cost improvements but showed none so I have not included them in their respective reports.

(1) Description of Improvement:

After I had completed my optimizations for all 5 tasks, I ran the following query:

Source: <https://docs.oracle.com/en/database/oracle/oracle-database/21/tqdba/tuning-database-buffer-cache.html#GUID-83733109-5119-4DDB-8A81-5302CE956BE2>

```
SQL> SELECT size_for_estimate, buffers_for_estimate, estd_physical_read_factor,
2          estd_physical_reads
3      FROM V$DB_CACHE_ADVICE
4     WHERE name = 'DEFAULT'
5          AND block_size = (SELECT value FROM V$PARAMETER WHERE name = 'db_block_size')
6          AND advice_status = 'ON';
```

Cache Size (MB)	Buffers	ESTD_PHYSICAL_READ_FACTOR	Estd Phys Reads
16	1,952	6.3196	1,953,987
32	3,904	5.2904	1,635,756
48	5,856	4.3002	1,329,581
64	7,808	3.4472	1,065,867
80	9,760	2.8281	874,444
96	11,712	2.291	708,361
112	13,664	1.8551	573,597
128	15,616	1.5287	472,669
144	17,568	1.2939	400,061
160	19,520	1.1366	351,432
176	21,472	1.0226	316,181
180	21,960	1	309,194
192	23,424	.9285	287,097
208	25,376	.8397	259,615

224	27,328	.7491	231,630
240	29,280	.6593	203,850
256	31,232	.5707	176,461
272	33,184	.4939	152,725
288	35,136	.4755	147,023
304	37,088	.4682	144,769
320	39,040	.4575	141,471

```
SQL> SHOW PARAMETER db_cache_size;
```

```
NAME          TYPE          VALUE
-----
db_cache_size big integer 184M
```

By default I have 184M already in the cache, meaning I can have up to a total of 184MB + 36MB = 222MB which is leftover from the given 100MB of transient storage space. The closest I can afford is 208MB which is estimated to reduce physical reads by just over 26%.

```
SQL> ALTER SYSTEM SET db_cache_size = 208M SCOPE=SPFILE;
```

```
System SET altered.
```

```
SQL> SHOW PARAMETER db_cache_size;
```

```
NAME          TYPE          VALUE
-----
db_cache_size big integer 208M
```

(2) Benefits of Improvement:

Allocating more transient storage to the data buffer cache significantly reduced the read operations as can be seen below

utlbstat/utlestat (Screenshots used because of spacing issues)

Before improvement:

TABLE_SPACE	FILE_NAME	READS	BLKS_READ	READ_TIME	WRITES	BLKS_WRT	WRITE_TIME	MEGABYTES	AVG_RT	blocks/rd
INDEX TS_32K	/opt/oracle/oradata/DB/32k_tbs.dbf	1	1	0	0	0	0	70	0	1
SYSAUX	/opt/oracle/oradata/DB/sysaux01.dbf	6	9	0	2	2	0	703	0	1.5
SYSTEM	/opt/oracle/oradata/DB/system01.dbf	50	55	0	48	99	1	965	0	1.1
TPCHR	/opt/oracle/oradata/DB/tpchr.dbf	206172	206199	37	0	0	0	3146	0	1
UNDOTBS1	/opt/oracle/oradata/DB/undotbs01.dbf	13	13	0	18	36	0	357	0	1
USERS	/opt/oracle/oradata/DB/users01.dbf	0	0	0	0	0	0	31	0	0

TPCHR READS = **206,172**

TPCHR BLKS_READ = **206,199**

After improvement:

TABLE_SPACE	FILE_NAME	READS	BLKS_READ	READ_TIME	WRITES	BLKS_WRT	WRITE_TIME	MEGABYTES	AVG_RT	blocks/rd
INDEX TS_32K	/opt/oracle/oradata/DB/32k_tbs.dbf	0	0	0	0	0	0	70	0	0
SYSAUX	/opt/oracle/oradata/DB/sysaux01.dbf	5	8	0	1	1	0	692	0	1.6
SYSTEM	/opt/oracle/oradata/DB/system01.dbf	14	14	0	57	106	2	954	0	1
TPCHR	/opt/oracle/oradata/DB/tpchr.dbf	16302	64731	6	0	0	0	3146	0	3.97
UNDOTBS1	/opt/oracle/oradata/DB/undotbs01.dbf	0	0	0	22	45	1	357	0	0
USERS	/opt/oracle/oradata/DB/users01.dbf	0	0	0	0	0	0	31	0	0

TPCHR READS = **16,302**

TPCHR BLKS_READ = **64,731**

Total READS improvement = 206,172 - 16,302 = **189,870**

Total BLKS_READ improvement = 206,199 - 64,731 = **141,468**

(3) Costs of Improvement:

The cost of allocating memory to the data buffer cache is 24MB of transient memory

```
SQL> SHOW PARAMETER db_cache_size;

NAME                TYPE          VALUE
-----
db_cache_size       big integer   208M
```

(4) Report from improvement:

SQL> CONNECT SYSTEM/oracle

Connected.

SQL> ALTER SYSTEM SET db_cache_size = 208M SCOPE=SPFILE;

System SET altered.

SQL> SHOW PARAMETER db_cache_size;

NAME	TYPE	VALUE

db_cache_size	big integer	208M

SQL>

SQL> CONNECT TPCHR/oracle

Connected.

SQL> SELECT S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE, SUM(TOTAL_REVENUE) AS TOTAL_REVENUE
2 FROM REVENUE
3 GROUP BY S_SUPPKEY, S_NAME, S_ADDRESS, S_PHONE
4 HAVING SUM(TOTAL_REVENUE) =
5 (SELECT MAX(SUM(R.TOTAL_REVENUE))
6 FROM REVENUE R
7 GROUP BY R.S_SUPPKEY)
8 ORDER BY S_SUPPKEY;

S_SUPPKEY	S_NAME	S_ADDRESS	S_PHONE	TOTAL_REVENUE

2993	Supplier#000002993	CDRN7azuEWTawl9G	30-541-514-5637	12267625.5

1 row selected.

CSCI317 – Report1 - Samuel Ian Black – SIB979 - 6025821

Total Costs

Persistent Storage: 7MB of 300MB

1. TASK1IDX1 = 6.5MB
2. TASK1IDX2 = 0.5MB

Transient Memory: 24MB of 100MB

1. db_cache_size = 208
 - Originally 184 + 24 from allocated 100MB Transient Storage expansion