

Table of Contents

Original cost of query.....	2
Optimization #1 – Transformation of Select Statement.....	3
Optimization #2 – Creation of Bitmap Index	9
Total Costs.....	15

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Original cost of query

N/A will not return result/s due to nested loops

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Optimization #1 – Transformation of Select Statement

Note that I decided against using Materialized Views as we're unsure of whether or not the client's DB has designated downtime to mitigate the more expensive insertion, and if the client doesn't have a designated downtime for the DB to perform inserts and commits to update the materialized views and instead wants the views to remain up to date and therefore, updated on insert, if the table LINEITEM is frequently inserted to which is what I assume since it's the largest table, the performance of the entire DB will greatly suffer which is why I've decided not to use Materialized Views.

(1) Description of Improvement:

The query was previously selecting the function TSIZE from the table LINEITEM

e.g. `SELECT DISTINCT TSIZE FROM LINEITEM` which is incorrect because the TSIZE function already queries the LINEITEM table therefore, this can be solved by changing the query to: `SELECT TSIZE FROM DUAL`.

(2) Benefits of Improvement:

The query will now end and return the correct result.

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE OR REPLACE FUNCTION TSIZE RETURN NUMBER
```

```
2 AS
```

```
3 TS NUMBER;
```

```
4 BEGIN
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

```
5    SELECT AVG(COUNT(*))
6    INTO TS
7    FROM LINEITEM
8    GROUP BY L_ORDERKEY;
9
10   RETURN(TS);
11  END;
12  /
```

Function TSIZE compiled

SQL>

SQL> show errors

SQL>

SQL> EXPLAIN PLAN FOR

```
2    SELECT COUNT(*)
3  FROM ( SELECT L_ORDERKEY, COUNT(*)
4         FROM LINEITEM
5         GROUP BY L_ORDERKEY
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

```
6          HAVING COUNT(*) > (SELECT TSIZE
7                                FROM DUAL) );
```

Explained.

```
SQL>
```

```
SQL> @showplan
```

```
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 4190422380

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		1		1621 (4)	00:00:01	
1	SORT AGGREGATE		1				
2	VIEW		1940K		1621 (4)	00:00:01	

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

*	3		FILTER									
	4		HASH GROUP BY				1940K	24M	1621	(4)	00:00:01	
	5		INDEX FAST FULL SCAN	LINEITEM_PKEY		1940K	24M	1571	(1)	00:00:01		
	6		FAST DUAL				1		2	(0)	00:00:01	

Predicate Information (identified by operation id):

3 - filter(COUNT(*) > (SELECT "TSIZE"() FROM "SYS"."DUAL" "DUAL"))

Total Cost before improvement = **N/A**

Total Cost after improvement = 1,621 + 1,621 + 1,621 + 1,571 + 2 = **6,436**

Total Cost improvement = **N/A**

(3) Costs of Improvement:

There is no cost associated with this improvement.

(4) Report from improvement:

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE OR REPLACE FUNCTION TSIZE RETURN NUMBER
```

```
2 AS
```

```
3 TS NUMBER;
```

```
4 BEGIN
```

```
5     SELECT AVG (COUNT (*))
```

```
6     INTO TS
```

```
7     FROM LINEITEM
```

```
8     GROUP BY L_ORDERKEY;
```

```
9
```

```
10     RETURN (TS) ;
```

```
11 END;
```

```
12 /
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Function TSIZE compiled

SQL>

SQL> show errors

SQL>

```
SQL> SELECT COUNT(*)
      2  FROM ( SELECT L_ORDERKEY, COUNT(*)
      3           FROM LINEITEM
      4           GROUP BY L_ORDERKEY
      5           HAVING COUNT(*) > (SELECT TSIZE
      6                                FROM DUAL) );
```

COUNT(*)

193031

1 row selected.

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Optimization #2 – Creation of Bitmap Index

(1) Description of Improvement:

I created a bitmap index over the `L_ORDERKEY` column;

The total number of rows in a table `LINEITEM` is 1,800,093

the total number of distinct rows of a column `L_ORDERKEY` is 450,000

meaning that the cardinality of `L_ORDERKEY` = $(450,000/1,800,093) * 100 = 24.99\%$ making the `L_ORDERKEY` a reasonable column for a bitmap to be created.

```
CREATE BITMAP INDEX TASK3IDX1 ON ORDERS (O_ORDERDATE) ;
```

(2) Benefits of Improvement:

The query now scans the Bitmap index `TASK3IDX1` rather than a primary key which is slightly faster.

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE BITMAP INDEX TASK3IDX1 ON LINEITEM (L_ORDERKEY) ;
```

```
INDEX TASK3IDX1 created.
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

SQL>

SQL> EXPLAIN PLAN FOR

```
2  SELECT COUNT(*)
3  FROM ( SELECT L_ORDERKEY, COUNT(*)
4          FROM LINEITEM
5          GROUP BY L_ORDERKEY
6          HAVING COUNT(*) > (SELECT TSIZE
7                               FROM DUAL) );
```

Explained.

SQL>

SQL> @showplan

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT

Plan hash value: 2915800778

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	

0	SELECT STATEMENT		1		1597 (4)	00:00:01	
1	SORT AGGREGATE		1				
2	VIEW		1940K		1597 (4)	00:00:01	
* 3	FILTER						
4	HASH GROUP BY		1940K	24M	1597 (4)	00:00:01	
5	BITMAP CONVERSION TO ROWIDS		1940K	24M	1547 (1)	00:00:01	
6	BITMAP INDEX FAST FULL SCAN	TASK3IDX1					
7	FAST DUAL		1		2 (0)	00:00:01	

Predicate Information (identified by operation id):

3 - filter(COUNT(*) > (SELECT "TSIZE"() FROM "SYS"."DUAL" "DUAL"))

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

Total Cost before improvement = **6,436**

Total Cost after improvement = $1,597 + 1,597 + 1,597 + 1,547 + 2 = \mathbf{6,340}$

Total Cost improvement = $6,436 - 6,340 = \mathbf{96}$

(3) Costs of Improvement:

The cost of creating the index is 13.5MB in persistent storage.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME';
```

```
old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME'
```

```
new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='TASK3IDX1'
```

```
Index Size (MB)
```

```
-----
```

```
13.5
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

(4) Report from improvement:

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE BITMAP INDEX TASK3IDX1 ON LINEITEM(L_ORDERKEY);
```

```
INDEX TASK3IDX1 created.
```

```
SQL>
```

```
SQL> SELECT COUNT(*)
```

```
2 FROM ( SELECT L_ORDERKEY, COUNT(*)
```

```
3        FROM LINEITEM
```

```
4        GROUP BY L_ORDERKEY
```

```
5        HAVING COUNT(*) > (SELECT TSIZE
```

```
6                               FROM DUAL) );
```

```
COUNT(*)
```

CSCI317 – Report3 - Samuel Ian Black – SIB979 - 6025821

193031

1 row selected.

Total Costs

Persistent Storage: 86.25MB of 300MB

1. TASK1IDX1 = 6.5MB
2. TASK1IDX2 = 0.5MB
3. TASK2IDX1 = 1.75MB
4. INDEX_TS_32K = 64MB (Size of Tablespace used to calculate Persistent Storage)
 - TASK2IDX2 = 60MB
5. TASK3IDX1 = 13.5MB

Transient Memory: 88MB of 100MB

1. db_32K_cache_size = 64M
2. db_cache_size = 208
 - Originally 184 + 24 from allocated 100MB Transient Storage expansion