

## Table of Contents

Original time of query:.....	2
Optimization #1 – Transformation of Select Statement.....	3
Optimization #2 – Creation of B*Tree Index .....	7
Time: .....	7
Explain Plan Before Creation of Index .....	9
Explain Plan After Creation of Index .....	11
Total Costs.....	17

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Note that before each run of a command time java task4 the below query is run.

```
CONNECT SYSTEM/oracle;
```

```
ALTER SYSTEM FLUSH SHARED_POOL;
```

Original time of query:

```
[oracle@localhost Exam]$ time java task4
```

```
Connected as tpchr user
```

```
Part key: 4471
```

```
Part key: 8526
```

```
Part key: 12224
```

```
Part key: 14988
```

```
Part key: 20143
```

```
Part key: 43902
```

```
Part key: 45138
```

```
Part key: 46172
```

```
Part key: 50506
```

```
Part key: 51788
```

```
Part key: 57573
```

```
Part key: 58525
```

```
Done.
```

```
real 0m1.307s
```

```
user 0m1.233s
```

```
sys 0m0.058s
```

Total time = 1.307 + 1.233 + 0.058 = **2.598s**

### Optimization #1 – Transformation of Select Statement

#### (1) Description of Improvement:

The results from the table BRAND55 are every row from the table PART where the column P\_BRAND = 'Brand#55'.

The results from the table NICKEL are every row from the table PART where the column P\_TYPE = 'ECONOMY BRUSHED NICKEL'.

The results from the table TEMP12 are every row from the table BRAND55, where P\_TYPE <> 'ECONOMY BRUSHED NICKEL'.

The results from the original query are every row from the table BRAND55 where P\_PARTKEY NOT IN the table TEMP12.

This means that the query is returning every row from the table PART where P\_BRAND = 'Brand#55' AND P\_TYPE = 'ECONOMY BRUSHED NICKEL'.

Therefore, the query can be simplified to the below without the need for creation of tables.

```
Statement stmt = conn.createStatement();
ResultSet rset = stmt.executeQuery(
    "SELECT P_PARTKEY " +
    "FROM PART " +
    "WHERE P_BRAND = 'Brand#55' " +
    "AND P_TYPE = 'ECONOMY BRUSHED NICKEL' " +
    "ORDER BY P_PARTKEY" );
```

#### (2) Benefits of Improvement:

1. No longer using persistent storage to create multiple relational tables
2. Improved query is much more readable
3. The improved query will also improve performance because of improvement 1.

```
[oracle@localhost Exam]$ time java task4
Connected as tpchr user
Part key: 4471
Part key: 8526
Part key: 12224
Part key: 14988
Part key: 20143
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Part key: 43902

Part key: 45138

Part key: 46172

Part key: 50506

Part key: 51788

Part key: 57573

Part key: 58525

Done.

real 0m0.839s

user 0m1.168s

sys 0m0.054s

real before improvement = **1.307s**

real improvement =  $1.307 - 0.839 = \mathbf{0.468s}$

user before improvement = **1.233s**

user improvement =  $1.233 - 1.168 = \mathbf{0.065s}$

sys before improvement = **0.058s**

sys improvement =  $0.058 - 0.054 = \mathbf{0.004s}$

Total time =  $0.839 + 1.168 + 0.054 = \mathbf{2.061s}$

Total time improvement =  $2.598 - 2.061 = \mathbf{0.537s}$

### (3) Costs of Improvement:

There is no cost increase associated with this improvement.

### (4) Report from improvement:

```
import java.sql.*;

class task4
{
    public static void main (String args [])
        throws SQLException, ClassNotFoundException
    {
        // Load the Oracle JDBC driver
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection
            ("jdbc:oracle:thin:@localhost:1521:db", "tpchr", "oracle");
        System.out.println( "Connected as tpchr user");
        try{

            Statement stmt = conn.createStatement();
            ResultSet rset = stmt.executeQuery(
                "SELECT P_PARTKEY " +
                "FROM PART " +
                "WHERE P_BRAND = 'Brand#55' " +
                "AND P_TYPE = 'ECONOMY BRUSHED NICKEL' " +
                "ORDER BY P_PARTKEY" );

            while ( rset.next() )
                System.out.println("Part key: " + rset.getInt(1) );

            System.out.println( "Done." );
        }
        catch (SQLException e )
        {
            String errmsg = e.getMessage();
            System.out.println( errmsg );
        }
    }
}
```

```
}  
}  
}
```

### Optimization #2 – Creation of B\*Tree Index

#### (1) Description of Improvement:

Creating an index `CREATE INDEX TASK4IDX1 ON PART(P_BRAND, P_TYPE, P_PARTKEY) COMPRESS 2;` allows for use of an index rather than accessing the entire `PART` table, compression also saves 2MB as opposed to not compressing and further decreases cost of the query since two prefix columns in the `TASK4IDX1` are repeated where as the primary key is unique and therefore, not.

#### (2) Benefits of Improvement:

The index allows for the query optimizer to traverse through the `TASK4IDX1` index vertically rather than accessing the entire `PART` table.

Time:

```
[oracle@localhost Exam]$ time java task4
```

```
Connected as tpchr user
```

```
Part key: 4471
```

```
Part key: 8526
```

```
Part key: 12224
```

```
Part key: 14988
```

```
Part key: 20143
```

```
Part key: 43902
```

```
Part key: 45138
```

```
Part key: 46172
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Part key: 50506

Part key: 51788

Part key: 57573

Part key: 58525

Done.

real 0m0.778s

user 0m1.118s

sys 0m0.054s

real before improvement = **0.839s**

real improvement =  $0.839 - 0.778 = \mathbf{0.061s}$

user before improvement = **1.168s**

user improvement =  $1.168 - 1.118 = \mathbf{0.05s}$

sys before improvement = **0.054s**

sys improvement =  $0.054 - 0.054 = \mathbf{0s}$

Total time =  $0.778 + 1.118 + 0.054 = \mathbf{1.95s}$

Total time improvement =  $2.061 - 1.95 = \mathbf{0.111s}$



## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Explain Plan Before Creation of Index

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> EXPLAIN PLAN FOR
```

```
2  SELECT P_PARTKEY
```

```
3  FROM PART
```

```
4  WHERE P_BRAND = 'Brand#55'
```

```
5  AND P_TYPE = 'ECONOMY BRUSHED NICKEL'
```

```
6  ORDER BY P_PARTKEY;
```

Explained.

```
SQL>
```

```
SQL> @showplan
```

```
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

-----  
Plan hash value: 2726178166

-----  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |  
-----  
0	SELECT STATEMENT		6	234	402 (1)	00:00:01
1	SORT ORDER BY		6	234	402 (1)	00:00:01
\* 2	TABLE ACCESS FULL	PART	6	234	401 (1)	00:00:01
-----

Predicate Information (identified by operation id):  
-----

2 - filter("P\_TYPE"='ECONOMY BRUSHED NICKEL' AND  
"P\_BRAND"='Brand#55')

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Explain Plan After Creation of Index

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE INDEX TASK4IDX1 ON PART(P_BRAND, P_TYPE, P_PARTKEY) COMPRESS 2;
```

Index TASK4IDX1 created.

```
SQL>
```

```
SQL> EXPLAIN PLAN FOR
```

```
2  SELECT P_PARTKEY
```

```
3  FROM PART
```

```
4  WHERE P_BRAND = 'Brand#55'
```

```
5  AND P_TYPE = 'ECONOMY BRUSHED NICKEL'
```

```
6  ORDER BY P_PARTKEY;
```

Explained.

```
SQL>
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

```
SQL> @showplan
```

```
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT
```

```
-----  
Plan hash value: 315755535
```

```
-----  
| Id  | Operation          | Name          | Rows  | Bytes | Cost (%CPU)| Time     |  
-----  
|  0  | SELECT STATEMENT   |               |    12 |   468 |    2   (0)| 00:00:01 |  
|*   1  |  INDEX RANGE SCAN | TASK4IDX1     |    12 |   468 |    2   (0)| 00:00:01 |  
-----
```

```
Predicate Information (identified by operation id):
```

```
-----  
  
1 - access("P_BRAND"='Brand#55' AND "P_TYPE"='ECONOMY BRUSHED  
      NICKEL')
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Total Cost before improvement =  $402 + 402 + 401 = \mathbf{1,205}$

Total Cost after improvement =  $2 + 2 = \mathbf{4}$

Total Cost improvement =  $1,205 - 4 = \mathbf{1201}$

### (3) Costs of Improvement:

The cost of creating the index is 1.5MB in persistent storage.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME';
```

```
old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME'
```

```
new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='TASK4IDX1'
```

Index Size (MB)

-----

1.5

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

### (4) Report from improvement:

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE INDEX TASK4IDX1 ON PART(P_BRAND, P_TYPE, P_PARTKEY) COMPRESS 2;
```

```
Index TASK4IDX1 created.
```

```
SQL>
```

```
SQL> SELECT P_PARTKEY
```

```
2 FROM PART
```

```
3 WHERE P_BRAND = 'Brand#55'
```

```
4 AND P_TYPE = 'ECONOMY BRUSHED NICKEL'
```

```
5 ORDER BY P_PARTKEY;
```

```
P_PARTKEY
```

```
-----
```

## CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

4471

8526

12224

14988

20143

43902

45138

46172

50506

51788

57573

58525

12 rows selected.





### Total Costs

Persistent Storage: 87.75MB of 300MB

1. TASK1IDX1 = 6.5MB
2. TASK1IDX2 = 0.5MB
3. TASK2IDX1 = 1.75MB
4. INDEX\_TS\_32K = 64MB (Size of Tablespace used to calculate Persistent Storage)
  - TASK2IDX2 = 60MB
5. TASK3IDX1 = 13.5MB
6. TASK4IDX1 = 1.5MB

Transient Memory: 88MB of 100MB

1. db\_32K\_cache\_size = 64M
2. db\_cache\_size = 208
  - Originally 184 + 24 from allocated 100MB Transient Storage expansion