

Table of Contents

Original time of query:.....	2
Optimization #1 – Transformation of Select Statement.....	4
Optimization #2 – Creation of B*Tree Index	10
Time	10
Explain Plan Before Creation of Index	11
Explain Plan After Creation of Index	12
Optimization #3 –Rebuild of TASK5IDX1 In INDEX_TS_32K Tablespace	18
Time	18
Explain Plan	19
utlbstat/utlestat (Screenshots used because of spacing issues)	20
Buffer Cache Hit Ratio.....	23
Total Costs.....	24

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Note that before each run of a command time java task5 the below query was run.

```
CONNECT SYSTEM/oracle;  
ALTER SYSTEM FLUSH SHARED_POOL;
```

Original time of query:

```
[oracle@localhost Exam]$ time java task5
```

```
Connected as CSCI317.
```

```
Total: 300
```

```
Done.
```

```
real 0m4.325s
```

```
user 0m3.720s
```

```
sys 0m0.890s
```

Total time = 4.325 + 3.720 + 0.890 = **8.935s**

Optimization #1 – Transformation of Select Statement

(1) Description of Improvement:

Removal of unnecessary outer while loop by using a prepared statement and transformation of select statement.

```
String query = "SELECT P_NAME " +  
               "FROM PART " +  
               "WHERE P_NAME = ?" ;  
  
PreparedStatement stmt = conn.prepareStatement (query);  
  
BufferedReader in = new BufferedReader(new  
FileReader("task5.txt"));  
  
    String str;  
    while ( ( str = in.readLine() ) != null ) {  
        stmt.setString(1, str);  
        ResultSet rset = stmt.executeQuery();  
        while (rset.next()) {  
            count ++;  
        }  
    }
```

(2) Benefits of Improvement:

The original solution had an extra loop which iterated through every single row of the table `PART`, and then a nested loop which iterated through the entire `task5.txt` file for each row in `PART`.

This was causing $m * n$ operations to be performed where m is the number of rows in the table `PART` and n is the number of lines in the `task5.txt` file.

The optimized solution performs n iterations where n is the number of lines in the `task5.txt` file and uses a prepared statement to query the database which avoids flooding the data buffer cache unnecessarily.

```
[oracle@localhost Exam]$ time java task5
```

```
Connected as CSCI317.
```

```
Total: 300
```

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Done.

real 0m1.560s

user 0m1.451s

sys 0m0.072s

real before improvement = **4.325s**

real improvement = $4.325 - 1.560 = \mathbf{2.765s}$

user before improvement = **3.720s**

user improvement = $3.720 - 1.451 = \mathbf{2.269s}$

sys before improvement = **0.890s**

sys improvement = $0.890 - 0.072 = \mathbf{0.818s}$

Total time = $1.560 + 1.451 + 0.072 = \mathbf{3.083s}$

Total time improvement = $8.935 - 3.083 = \mathbf{5.852s}$

(3) Costs of Improvement:

There is no cost associated with this improvement.

(4) Report from improvement:

```
import java.sql.*;
import java.io.*;

class task5
{
    public static void main (String args [])
        throws SQLException, ClassNotFoundException
    {
        // Load the Oracle JDBC driver
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection
            ("jdbc:oracle:thin:@localhost:1521:db", "tpchr", "oracle");
        System.out.println( "Connected as CSCI317." );
        try{
            int count = 0;

            String query = "SELECT P_NAME " +
                           "FROM PART " +
                           "WHERE P_NAME = ?" ;
```

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

```
PreparedStatement stmt = conn.prepareStatement (query);

BufferedReader in = new BufferedReader( new FileReader("task5.txt") );
String str;
while ( ( str = in.readLine() ) != null )
{
    stmt.setString(1, str);
    ResultSet rset = stmt.executeQuery();
    while (rset.next()) {
        count ++;
    }
}
in.close();

System.out.println( "Total: " + count );
System.out.println( "Done." );
}
catch (SQLException e)
{

```


CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

```
String errmsg = e.getMessage();  
System.out.println( errmsg );  
}  
catch (IOException io )  
{  
    String errmsg = io.getMessage();  
    System.out.println( errmsg );  
}  
  
}  
}
```

Optimization #2 – Creation of B*Tree Index

(1) Description of Improvement:

Creation of an index on P_NAME .

```
CREATE INDEX TASK5IDX1 ON PART(P_NAME) ;
```

(2) Benefits of Improvement:

The creation of the index allows for the query optimizer to traverse the index vertically rather than having to access the entire table.

Time

```
[oracle@localhost Exam]$ time java task5
```

```
Connected as CSCI317.
```

```
Total: 300
```

```
Done.
```

```
real 0m0.956s
```

```
user 0m1.348s
```

```
sys 0m0.063s
```

real before improvement = 1.560s

real improvement = $1.560 - 0.956 = \mathbf{0.604s}$

user before improvement = 1.451s

user improvement = $1.451 - 1.348 = \mathbf{0.103s}$

sys before improvement = 0.072s

sys improvement = $0.072 - 0.063 = \mathbf{0.009s}$

Total time = $0.956 + 1.348 + 0.063 = \mathbf{2.367s}$

Total time improvement = $3.083 - 2.367 = \mathbf{0.716s}$

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Explain Plan Before Creation of Index

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> EXPLAIN PLAN FOR
```

```
2  SELECT P_NAME
```

```
3  FROM PART
```

```
4  WHERE P_NAME = 'turquoise yellow magenta burnished peach';
```

Explained.

```
SQL>
```

```
SQL> @showplan
```

```
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 673417232

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	

0	SELECT STATEMENT		12	348	401 (1)	00:00:01	
* 1	TABLE ACCESS FULL	PART	12	348	401 (1)	00:00:01	

Predicate Information (identified by operation id):

1 - filter("P_NAME"='turquoise yellow magenta burnished peach')

Explain Plan After Creation of Index

SQL> SET ECHO ON

SQL> SET FEEDBACK ON

SQL> SET LINESIZE 300

SQL> SET PAGESIZE 300

SQL>

SQL> CREATE INDEX TASK5IDX1 ON PART(P_NAME);

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Index TASK5IDX1 created.

SQL>

SQL> EXPLAIN PLAN FOR

2 SELECT P_NAME

3 FROM PART

4 WHERE P_NAME = 'turquoise yellow magenta burnished peach';

Explained.

SQL>

SQL> @showplan

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

PLAN_TABLE_OUTPUT

Plan hash value: 984302744

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	

0	SELECT STATEMENT		1	29	3 (0)	00:00:01	
* 1	INDEX RANGE SCAN	TASK5IDX1	1	29	3 (0)	00:00:01	

Predicate Information (identified by operation id):

1 - access("P_NAME"='turquoise yellow magenta burnished peach')

Total Cost before improvement = 401 + 401 = **802**

Total Cost after improvement = 3 + 3 = **6**

Total Cost improvement = 802 – 6 = **796**

(3) Costs of Improvement:

The cost of creating the index is 3.25MB in persistent storage.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME';
```

```
old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='&INDEX_NAME'
```

```
new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where  
segment_name='TASK5IDX1'
```

```
Index Size (MB)
```

```
-----
```

```
3.25
```

(4) Report from improvement:

```
SQL> SET ECHO ON
```

```
SQL> SET FEEDBACK ON
```

```
SQL> SET LINESIZE 300
```

```
SQL> SET PAGESIZE 300
```

```
SQL>
```

```
SQL> CREATE INDEX TASK5IDX1 ON PART(P_NAME);
```

```
Index TASK5IDX1 created.
```

```
SQL>
```

```
SQL> SELECT P_NAME
```

```
2 FROM PART
```

```
3 WHERE P_NAME = 'turquoise yellow magenta burnished peach';
```

```
P_NAME
```

```
-----
```

```
turquoise yellow magenta burnished peach
```


CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

1 row selected.

CSCI317 – Report4 - Samuel Ian Black – SIB979 - 6025821

Optimization #3 –Rebuild of TASK5IDX1 In INDEX_TS_32K Tablespace

(1) Description of Improvement:

Because TASK5IDX1 is a B*Tree index, I can take full advantage of storing it in the 32k data buffer cache which will better balance the tree and reduce physical and logical reads of the query due to “flattening” the tree over larger data blocks.

```
ALTER INDEX TASK5IDX1 REBUILD TABLESPACE INDEX_TS_32K;
```

(2) Benefits of Improvement:

The B*Tree is now more balanced which results in significantly less data blocks needed to be read resulting in less physical reads and reduced Cost of performing the query.

Time

```
[oracle@localhost Exam]$ time java task5
```

```
Connected as CSCI317.
```

```
Total: 300
```

```
Done.
```

```
real 0m0.899s
```

```
user 0m1.306s
```

```
sys 0m0.076s
```

```
real before improvement = 0.956s
```

```
real improvement = 0.956 – 0.899 = 0.057s
```

```
user before improvement = 1.348s
```

```
user improvement = 1.348 – 1.306 = 0.042s
```

```
sys before improvement = 0.063s
```

```
sys improvement = 0.063 – 0.076 = -0.013s
```

```
Total time = 0.899 + 1.306 + -0.013 = 2.192s
```

```
Total time improvement = 2.367 – 2.192 = 0.175s
```

Explain Plan

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> ALTER INDEX TASK5IDX1 REBUILD TABLESPACE INDEX_TS_32K;
```

Index TASK5IDX1 altered.

```
SQL>
SQL> EXPLAIN PLAN FOR
  2  SELECT P_NAME
  3  FROM PART
  4  WHERE P_NAME = 'turquoise yellow magenta burnished peach';
```

Explained.

```
SQL>
SQL> @showplan
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 984302744

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time

	0		SELECT STATEMENT				1		29		1	(0)		00:00:01	
	*		INDEX RANGE SCAN		TASK5IDX1		1		29		1	(0)		00:00:01	

Predicate Information (identified by operation id):

1 - access("P_NAME"='turquoise yellow magenta burnished peach')

Total Cost before improvement = 3 + 3 = 6

Total Cost after improvement = 1 + 1 = 2

Total Cost improvement = 6 – 2 = 4

utlbstat/utlestat (Screenshots used because of spacing issues)

Before improvement:

TABLE_SPACE	FILE_NAME	READS	BLKS_READ	READ_TIME	WRITES	BLKS_WRT	WRITE_TIME	MEGABYTES	AVG_RT	blocks/rd
INDEX_TS_32K	/opt/oracle/oradata/DB/32k_tbs.dbf	0	0	0	0	0	0	67	0	0
SYSAUX	/opt/oracle/oradata/DB/sysaux01.dbf	0	0	0	0	0	0	692	0	0
SYSTEM	/opt/oracle/oradata/DB/system01.dbf	19	19	0	0	0	0	954	0	1
TPCHR	/opt/oracle/oradata/DB/tpchr.dbf	30	277	0	0	0	0	3146	0	9.23
UNDOTBS1	/opt/oracle/oradata/DB/undotbs01.dbf	0	0	0	0	0	0	357	0	0
USERS	/opt/oracle/oradata/DB/users01.dbf	0	0	0	0	0	0	31	0	0

TPCHR + INDEX_TS_32K READS = 30

TPCHR + INDEX_TS_32K BLKS_READ = 277

After improvement:

TABLE_SPACE	FILE_NAME	READS	BLKS_READ	READ_TIME	WRITES	BLKS_WRT	WRITE_TIME	MEGABYTES	AVG_RT	blocks/rd
INDEX_TS_32K	/opt/oracle/oradata/DB/32k_tbs.dbf	30	33	0	0	0	0	67	0	1.1
SYSAUX	/opt/oracle/oradata/DB/sysaux01.dbf	4	7	0	0	0	0	692	0	1.75
SYSTEM	/opt/oracle/oradata/DB/system01.dbf	18	18	0	0	0	0	954	0	1
TPCHR	/opt/oracle/oradata/DB/tpchr.dbf	0	0	0	0	0	0	3146	0	0
UNDOTBS1	/opt/oracle/oradata/DB/undotbs01.dbf	0	0	0	0	0	0	357	0	0
USERS	/opt/oracle/oradata/DB/users01.dbf	0	0	0	0	0	0	31	0	0

TPCHR + INDEX_TS_32K READS = 30

TPCHR + INDEX_TS_32K BLKS_READ = 33

Total READS improvement = 30 – 30 = **0**

Total BLKS_READ improvement = 277 – 33 = **244**

(3) Costs of Improvement:

Rebuilding of the index TASK5IDX1 does not add additional persistent storage, the index is simply rebuilt on the 32K Tablespace.

```
SQL> select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where segment_name='&INDEX_NAME';
old:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where segment_name='&INDEX_NAME'
new:select sum(bytes)/1024/1024 as "Index Size (MB)" from dba_segments where segment_name='TASK5IDX1'
```

Index Size (MB)

3.25

(4) Report from improvement:

```
SQL> SET ECHO ON
SQL> SET FEEDBACK ON
SQL> SET LINESIZE 300
SQL> SET PAGESIZE 300
SQL>
SQL> ALTER INDEX TASK5IDX1 REBUILD TABLESPACE INDEX_TS_32K;
```

Index TASK5IDX1 altered.

```
SQL>
SQL> SELECT P_NAME
2 FROM PART
3 WHERE P_NAME = 'turquoise yellow magenta burnished peach';
```

P_NAME

turquoise yellow magenta burnished peach

1 row selected.

Buffer Cache Hit Ratio

Because I’ve been altering transient storage, It’s important that I ensure the hit rate is at least 85% but preferably above 95%.

Source:[https://docs.oracle.com/database/121/TGDBA/tune_buffer_cache.htm#TGDBA536]

```
SQL> SELECT name, value
2  FROM V$SYSSTAT
3  WHERE name IN ('db block gets from cache', 'consistent gets from cache',
4  'physical reads cache');
```

NAME	VALUE
-----	-----
db block gets from cache	23137
consistent gets from cache	6889631
physical reads cache	742706

Buffer Cache Hit Ratio = 1 - (('physical reads cache') / ('consistent gets from cache' + 'db block gets from cache'))

= 1 – ((742706) / (6889631+ 23137)) = 0.8925602595 or 89.25%

This result was obtained after running every query and every java script twice, indicating that after normal use will stabilize at a higher hit rate which is an indicator that the buffer cache has been used appropriately to optimize the Database.

Total Costs

Persistent Storage: 87.75MB of 300MB

1. TASK1IDX1 = 6.5MB
2. TASK1IDX2 = 0.5MB
3. TASK2IDX1 = 1.75MB
4. INDEX_TS_32K = 64MB (Size of Tablespace used to calculate Persistent Storage)
 - TASK2IDX2 = 60MB
 - TASK5IDX1 3.25MB
5. TASK3IDX1 = 13.5MB
6. TASK4IDX1 = 1.5MB

Transient Memory: 88MB of 100MB

1. db_32K_cache_size = 64M
2. db_cache_size = 208
 - Originally 184 + 24 from allocated 100MB Transient Storage expansion